

BCG_Model

June 15, 2021

```
[5]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import datetime
```

1 Main Direction

- Sub-task 1: Think through what key drivers of churn could be for our client
- Sub-task 2: Build the features in order to get ready to model

From previous EDA,

Firstly, I found that the churn percentage is about 10%. And we saw that churn is not specifically related to any SME category in particular. As for sales channel, some channel has 0% churn rate. However it may be due to the fact that there are only a few customers through certain channels.

Secondly, I reckoned the consumption-related factors (cons_12m, cons_fas_12m ...) may contribute to this churn percentage. We can observe a highly right-skewed distribution in most of the consumption variables. From the box-plots we can discover there are plenty of outliers, which further indicate the skewness of consumptions.

The “forecast of consumption” can contribute substantially as well, however the large missing value (>75%) makes some of these data unable to provide useful information, if we fill up with estimate values, it adds confusion to the model.

Further, for the attributes such as “num_years_antig”, we can clearly see that less loyal customers (low years of antig) are more likely to churn. And customers with many active products and services are less likely to churn.

Lastly, I observed that at certain periods of time, the churn rate is significantly higher, which may result from some economic cycle or some impact on the market. So I believe this may contribute useful information to our model.

Potential key drivers:

1. Consumptions.
2. num_years_antig, nb_prod_act ...etc
3. Dates

2 Feature Engineering

```
[6]: train_data = pd.read_csv('ml_case_training_data.csv')
      churn_data = pd.read_csv('ml_case_training_output.csv')
      history_data = pd.read_csv('ml_case_training_hist_data.csv')
```

```
[7]: train_data.head()
```

```
[7]:
```

	id	activity_new	\
0	48ada52261e7cf58715202705a0451c9	esoiifxdlbkcsluxmfuacbdckommixw	
1	24011ae4ebbe3035111d65fa7c15bc57		NaN
2	d29c2c54acc38ff3c0614d0a653813dd		NaN
3	764c75f661154dac3a6c254cd082ea7d		NaN
4	bba03439a292a1e166f80264c16191cb		NaN

	campaign_disc_ele	channel_sales	cons_12m	\
0	NaN	lmkebamcaaclubfxadlmueccxoimlema	309275	
1	NaN	foosdfpfkusacimwkcsosbicdxkicaua	0	
2	NaN	NaN	4660	
3	NaN	foosdfpfkusacimwkcsosbicdxkicaua	544	
4	NaN	lmkebamcaaclubfxadlmueccxoimlema	1584	

	cons_gas_12m	cons_last_month	date_activ	date_end	date_first_activ	\
0	0	10025	2012-11-07	2016-11-06		NaN
1	54946	0	2013-06-15	2016-06-15		NaN
2	0	0	2009-08-21	2016-08-30		NaN
3	0	0	2010-04-16	2016-04-16		NaN
4	0	0	2010-03-30	2016-03-30		NaN

	... forecast_price_pow_p1	has_gas	imp_cons	margin_gross_pow_ele	\
0	...	58.995952	f	831.8	-41.76
1	...	40.606701	t	0.0	25.44
2	...	44.311378	f	0.0	16.38
3	...	44.311378	f	0.0	28.60
4	...	44.311378	f	0.0	30.22

	margin_net_pow_ele	nb_prod_act	net_margin	num_years_antig	\
0	-41.76	1	1732.36	3	
1	25.44	2	678.99	3	
2	16.38	1	18.89	6	
3	28.60	1	6.60	6	
4	30.22	1	25.46	6	

	origin_up	pow_max
0	ldkssxwpmemidmecebumciepifcamkci	180.000
1	lxidpiddsbxsbosboudacockeimpuepw	43.648
2	kamkkxfxxuwbdslkwifmmcsiusiusws	13.800

```
3 kamkkxfxxuwbdslkwifmmcsiusiosws 13.856
4 kamkkxfxxuwbdslkwifmmcsiusiosws 13.200
```

[5 rows x 32 columns]

```
[8]: train_data.describe()
```

```
[8]:
```

	campaign_disc_ele	cons_12m	cons_gas_12m	cons_last_month	\
count	0.0	1.609600e+04	1.609600e+04	1.609600e+04	
mean	NaN	1.948044e+05	3.191164e+04	1.946154e+04	
std	NaN	6.795151e+05	1.775885e+05	8.235676e+04	
min	NaN	-1.252760e+05	-3.037000e+03	-9.138600e+04	
25%	NaN	5.906250e+03	0.000000e+00	0.000000e+00	
50%	NaN	1.533250e+04	0.000000e+00	9.010000e+02	
75%	NaN	5.022150e+04	0.000000e+00	4.127000e+03	
max	NaN	1.609711e+07	4.188440e+06	4.538720e+06	

	forecast_base_bill_ele	forecast_base_bill_year	forecast_bill_12m	\
count	3508.000000	3508.000000	3508.000000	
mean	335.843857	335.843857	3837.441866	
std	649.406000	649.406000	5425.744327	
min	-364.940000	-364.940000	-2503.480000	
25%	0.000000	0.000000	1158.175000	
50%	162.955000	162.955000	2187.230000	
75%	396.185000	396.185000	4246.555000	
max	12566.080000	12566.080000	81122.630000	

	forecast_cons	forecast_cons_12m	forecast_cons_year	...	\
count	3508.000000	16096.000000	16096.000000	...	
mean	206.845165	2370.555949	1907.347229	...	
std	455.634288	4035.085664	5257.364759	...	
min	0.000000	-16689.260000	-85627.000000	...	
25%	0.000000	513.230000	0.000000	...	
50%	42.215000	1179.160000	378.000000	...	
75%	228.117500	2692.077500	1994.250000	...	
max	9682.890000	103801.930000	175375.000000	...	

	forecast_price_energy_p1	forecast_price_energy_p2	\
count	15970.000000	15970.000000	
mean	0.135901	0.052951	
std	0.026252	0.048617	
min	0.000000	0.000000	
25%	0.115237	0.000000	
50%	0.142881	0.086163	
75%	0.146348	0.098837	
max	0.273963	0.195975	

	forecast_price_pow_p1	imp_cons	margin_gross_pow_ele	\
count	15970.000000	16096.000000	16083.000000	
mean	43.533496	196.123447	22.462276	
std	5.212252	494.366979	23.700883	
min	-0.122184	-9038.210000	-525.540000	
25%	40.606701	0.000000	11.960000	
50%	44.311378	44.465000	21.090000	
75%	44.311378	218.090000	29.640000	
max	59.444710	15042.790000	374.640000	

	margin_net_pow_ele	nb_prod_act	net_margin	num_years_antig	\
count	16083.000000	16096.000000	16081.000000	16096.000000	
mean	21.460318	1.347788	217.987028	5.030629	
std	27.917349	1.459808	366.742030	1.676101	
min	-615.660000	1.000000	-4148.990000	1.000000	
25%	11.950000	1.000000	51.970000	4.000000	
50%	20.970000	1.000000	119.680000	5.000000	
75%	29.640000	1.000000	275.810000	6.000000	
max	374.640000	32.000000	24570.650000	16.000000	

	pow_max
count	16093.000000
mean	20.604131
std	21.772421
min	1.000000
25%	12.500000
50%	13.856000
75%	19.800000
max	500.000000

[8 rows x 22 columns]

```
[9]: train = train_data.merge(churn_data, on = 'id')
```

2.1 Finding null and discard useless columns

```
[10]: null_percentage = (train_data.isnull().sum() / train_data.isnull().count()).
      ↪ sort_values(ascending = False)
missing_data = pd.DataFrame({'Total Null': train_data.isnull().sum(),
      'Null Percentage': null_percentage
      }).sort_values(by='Null Percentage', ascending_
      ↪ = False)
```

```
[11]: missing_data
```

```
[11]:
```

	Total Null	Null Percentage
campaign_disc_ele	16096	1.000000
date_first_activ	12588	0.782058
forecast_cons	12588	0.782058
forecast_bill_12m	12588	0.782058
forecast_base_bill_year	12588	0.782058
forecast_base_bill_ele	12588	0.782058
activity_new	9545	0.593004
channel_sales	4218	0.262053
date_modif_prod	157	0.009754
forecast_price_energy_p1	126	0.007828
forecast_price_pow_p1	126	0.007828
forecast_price_energy_p2	126	0.007828
forecast_discount_energy	126	0.007828
origin_up	87	0.005405
date_renewal	40	0.002485
net_margin	15	0.000932
margin_gross_pow_ele	13	0.000808
margin_net_pow_ele	13	0.000808
pow_max	3	0.000186
date_end	2	0.000124
forecast_meter_rent_12m	0	0.000000
forecast_cons_12m	0	0.000000
has_gas	0	0.000000
id	0	0.000000
imp_cons	0	0.000000
date_activ	0	0.000000
cons_last_month	0	0.000000
nb_prod_act	0	0.000000
cons_gas_12m	0	0.000000
num_years_antig	0	0.000000
cons_12m	0	0.000000
forecast_cons_year	0	0.000000

```
[12]: discard_col = missing_data[missing_data['Null Percentage'] > 0.5].index
```

```
[13]: discard_col
```

```
[13]: Index(['campaign_disc_ele', 'date_first_activ', 'forecast_cons',
            'forecast_bill_12m', 'forecast_base_bill_year',
            'forecast_base_bill_ele', 'activity_new'],
            dtype='object')
```

```
[14]: train = train.drop(columns = discard_col, axis = 0)
```

```
[15]: train.shape
```

```
[15]: (16096, 26)
```

```
[16]: num_col = train._get_numeric_data().columns.tolist()
      cat_col = set(train.columns) - set(num_col)
```

```
[17]: num_col
```

```
[17]: ['cons_12m',
      'cons_gas_12m',
      'cons_last_month',
      'forecast_cons_12m',
      'forecast_cons_year',
      'forecast_discount_energy',
      'forecast_meter_rent_12m',
      'forecast_price_energy_p1',
      'forecast_price_energy_p2',
      'forecast_price_pow_p1',
      'imp_cons',
      'margin_gross_pow_ele',
      'margin_net_pow_ele',
      'nb_prod_act',
      'net_margin',
      'num_years_antig',
      'pow_max',
      'churn']
```

```
[18]: cat_col
```

```
[18]: {'channel_sales',
      'date_activ',
      'date_end',
      'date_modif_prod',
      'date_renewal',
      'has_gas',
      'id',
      'origin_up'}
```

```
[19]: train.isnull().sum()
```

```
[19]: id                0
      channel_sales    4218
      cons_12m         0
      cons_gas_12m     0
      cons_last_month  0
      date_activ       0
      date_end         2
      date_modif_prod  157
```

```

date_renewal          40
forecast_cons_12m      0
forecast_cons_year     0
forecast_discount_energy 126
forecast_meter_rent_12m  0
forecast_price_energy_p1 126
forecast_price_energy_p2 126
forecast_price_pow_p1   126
has_gas               0
imp_cons              0
margin_gross_pow_ele   13
margin_net_pow_ele     13
nb_prod_act           0
net_margin            15
num_years_antig        0
origin_up             87
pow_max               3
churn                 0
dtype: int64

```

2.2 Deal with consumption

```
[20]: consumption_col = ['cons_12m', 'cons_gas_12m', 'cons_last_month']
      consumption = train[consumption_col]
```

```
[21]: consumption.describe()
```

```
[21]:
```

	cons_12m	cons_gas_12m	cons_last_month
count	1.609600e+04	1.609600e+04	1.609600e+04
mean	1.948044e+05	3.191164e+04	1.946154e+04
std	6.795151e+05	1.775885e+05	8.235676e+04
min	-1.252760e+05	-3.037000e+03	-9.138600e+04
25%	5.906250e+03	0.000000e+00	0.000000e+00
50%	1.533250e+04	0.000000e+00	9.010000e+02
75%	5.022150e+04	0.000000e+00	4.127000e+03
max	1.609711e+07	4.188440e+06	4.538720e+06

Firstly, the minimum of consumption is negative, which seems not plausible, may conclude this is due to corrupted data.

```
[22]: for col in consumption_col:
      consumption.loc[consumption[col] < 0, col] = 0
      # Check whether negative value left
      # print(consumption[consumption[col] < 0])
```

/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1637:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_block(indexer, value, name)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:692:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iloc._setitem_with_indexer(indexer, value, self.name)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1637:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_block(indexer, value, name)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:692:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iloc._setitem_with_indexer(indexer, value, self.name)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1637:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_block(indexer, value, name)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:692:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iloc._setitem_with_indexer(indexer, value, self.name)
```

```
[23]: consumption.describe()
```

```
[23]:
```

	cons_12m	cons_gas_12m	cons_last_month
count	1.609600e+04	1.609600e+04	1.609600e+04
mean	1.948339e+05	3.191227e+04	1.953172e+04
std	6.795045e+05	1.775883e+05	8.231606e+04

min	0.000000e+00	0.000000e+00	0.000000e+00
25%	5.906250e+03	0.000000e+00	0.000000e+00
50%	1.533250e+04	0.000000e+00	9.010000e+02
75%	5.022150e+04	0.000000e+00	4.127000e+03
max	1.609711e+07	4.188440e+06	4.538720e+06

```
[24]: def capping_outlier(df,col):
    q1 = df[col].quantile(0.25)
    q3 = df[col].quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - (1.5 * iqr)
    upper_bound = q3 + (1.5 * iqr)
    df.loc[df[col] < lower_bound, col ] = lower_bound
    df.loc[df[col] > upper_bound, col ] = upper_bound
    # detect_outlier(df,col)
    return df[col]

def detect_outlier(df, col):
    q1 = df[col].quantile(0.25)
    q3 = df[col].quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - (1.5 * iqr)
    upper_bound = q3 + (1.5 * iqr)
    l_outlier = df[col].apply(lambda x: x < lower_bound).sum()
    u_outlier = df[col].apply(lambda x: x > upper_bound).sum()
    total = df[col].count()

    print("Q1:{} Q3:{}".format(q1,q3))
    print("lower outlier :{}, Upper outliers: {}".format(l_outlier/
    ↳total,u_outlier/total))
```

```
[25]: for col in consumption_col:
    print(col)
    detect_outlier(consumption, col)
```

```
cons_12m
Q1:5906.25 Q3:50221.5
lower outlier :0.0, Upper outliers: 0.15780318091451292
cons_gas_12m
Q1:0.0 Q3:0.0
lower outlier :0.0, Upper outliers: 0.18141153081510936
cons_last_month
Q1:0.0 Q3:4127.0
lower outlier :0.0, Upper outliers: 0.15339214711729623
```

2.2.1 Depend on the model we choose, such model may be non-sensitive to outliers.

Further, majority of 'cons_gas_12m' is zero, almost all of its elements are outliers, if we capping these data may not able to truly reflect. **Not capping cons_gas_12m**

```
[26]: capping_col = set(consumption_col)-set(['cons_gas_12m'])

for col in capping_col:
    print(col)
    consumption[col] = capping_outlier(consumption, col)
```

```
cons_12m
```

```
cons_last_month
```

```
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
self._setitem_single_column(loc, value, pi)
```

```
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
self._setitem_single_column(loc, value, pi)
```

```
<ipython-input-26-4b62715a02ed>:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
consumption[col] = capping_outlier(consumption, col)
```

```
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
self._setitem_single_column(loc, value, pi)
```

```
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(loc, value, pi)
<ipython-input-26-4b62715a02ed>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
consumption[col] = capping_outlier(consumption, col)
```

```
[27]: consumption.describe()
```

```
[27]:
```

	cons_12m	cons_gas_12m	cons_last_month
count	16096.000000	1.609600e+04	16096.000000
mean	35514.824273	3.191227e+04	2802.216793
std	40977.732430	1.775883e+05	3741.789990
min	0.000000	0.000000e+00	0.000000
25%	5906.250000	0.000000e+00	0.000000
50%	15332.500000	0.000000e+00	901.000000
75%	50221.500000	0.000000e+00	4127.000000
max	116694.375000	4.188440e+06	10317.500000

```
[28]: for col in consumption_col:
      print(col)
      detect_outlier(consumption, col)
```

```
cons_12m
Q1:5906.25 Q3:50221.5
lower outlier :0.0, Upper outliers: 0.0
cons_gas_12m
Q1:0.0 Q3:0.0
lower outlier :0.0, Upper outliers: 0.18141153081510936
cons_last_month
Q1:0.0 Q3:4127.0
lower outlier :0.0, Upper outliers: 0.0
```

```
[29]: consumption.isnull().sum()
```

```
[29]: cons_12m          0
      cons_gas_12m    0
      cons_last_month 0
      dtype: int64
```

2.3 Deal with forecast

```
[30]: forecast_col = ['forecast_cons_12m', 'forecast_cons_year',  
    ↪ 'forecast_discount_energy', 'forecast_meter_rent_12m', 'forecast_price_energy_p1', 'forecast_p  
forecast = train[forecast_col]
```

```
[31]: forecast.describe()
```

```
[31]:
```

	forecast_cons_12m	forecast_cons_year	forecast_discount_energy \
count	16096.000000	16096.000000	15970.000000
mean	2370.555949	1907.347229	0.991547
std	4035.085664	5257.364759	5.160969
min	-16689.260000	-85627.000000	0.000000
25%	513.230000	0.000000	0.000000
50%	1179.160000	378.000000	0.000000
75%	2692.077500	1994.250000	0.000000
max	103801.930000	175375.000000	50.000000

	forecast_meter_rent_12m	forecast_price_energy_p1 \
count	16096.000000	15970.000000
mean	70.309945	0.135901
std	79.023251	0.026252
min	-242.960000	0.000000
25%	16.230000	0.115237
50%	19.440000	0.142881
75%	131.470000	0.146348
max	2411.690000	0.273963

	forecast_price_energy_p2	forecast_price_pow_p1
count	15970.000000	15970.000000
mean	0.052951	43.533496
std	0.048617	5.212252
min	0.000000	-0.122184
25%	0.000000	40.606701
50%	0.086163	44.311378
75%	0.098837	44.311378
max	0.195975	59.444710

```
[32]: for col in forecast_col:  
    forecast.loc[forecast[col]< 0, col] = 0  
#     Check whther negative value left  
#     print(forecast[forecast[col] < 0])
```

/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
`self._setitem_single_column(loc, value, pi)`

```
[33]: forecast.describe()
```

```
[33]:
```

	forecast_cons_12m	forecast_cons_year	forecast_discount_energy \
count	16096.000000	16096.000000	15970.000000
mean	2375.199540	1918.633387	0.991547
std	4027.190618	5204.270838	5.160969
min	0.000000	0.000000	0.000000
25%	513.230000	0.000000	0.000000
50%	1179.160000	378.000000	0.000000
75%	2692.077500	1994.250000	0.000000
max	103801.930000	175375.000000	50.000000

	forecast_meter_rent_12m	forecast_price_energy_p1 \
count	16096.000000	15970.000000
mean	70.333054	0.135901
std	78.974201	0.026252
min	0.000000	0.000000
25%	16.230000	0.115237
50%	19.440000	0.142881
75%	131.470000	0.146348
max	2411.690000	0.273963

	forecast_price_energy_p2	forecast_price_pow_p1
count	15970.000000	15970.000000
mean	0.052951	43.533503
std	0.048617	5.212188
min	0.000000	0.000000
25%	0.000000	40.606701
50%	0.086163	44.311378
75%	0.098837	44.311378
max	0.195975	59.444710

```
[34]: forecast.isnull().sum()
```

```
[34]: forecast_cons_12m      0
forecast_cons_year      0
forecast_discount_energy  126
forecast_meter_rent_12m   0
forecast_price_energy_p1  126
forecast_price_energy_p2  126
forecast_price_pow_p1     126
dtype: int64
```

```
[35]: for col in forecast_col:
      # fill null with median
      forecast[col].fillna(forecast[col].median(), inplace = True)
```

/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/series.py:4460:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    return super().fillna(
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/series.py:4460:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    return super().fillna(
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/series.py:4460:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    return super().fillna(
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/series.py:4460:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    return super().fillna(
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/series.py:4460:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    return super().fillna(
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/series.py:4460:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    return super().fillna(
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/series.py:4460:
```

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
return super().fillna(
```

```
[36]: for col in forecast_col:
       print(col)
       detect_outlier(forecast, col)
```

```
forecast_cons_12m
Q1:513.23 Q3:2692.07750000000003
lower outlier :0.0, Upper outliers: 0.08505218687872763
forecast_cons_year
Q1:0.0 Q3:1994.25
lower outlier :0.0, Upper outliers: 0.09903081510934393
forecast_discount_energy
Q1:0.0 Q3:0.0
lower outlier :0.0, Upper outliers: 0.03597166998011928
forecast_meter_rent_12m
Q1:16.23 Q3:131.47
lower outlier :0.0, Upper outliers: 0.02379473161033797
forecast_price_energy_p1
Q1:0.11523699999999999 Q3:0.146348
lower outlier :0.00621272365805169, Upper outliers: 0.0228006958250497
forecast_price_energy_p2
Q1:0.0 Q3:0.098837
lower outlier :0.0, Upper outliers: 0.0
forecast_price_pow_p1
Q1:40.606701 Q3:44.31137796
lower outlier :0.006336978131212724, Upper outliers: 0.04653330019880716
```

```
[37]: capping_col = set(forecast_col)-set(['forecast_discount_energy'])
       print(capping_col)
       for col in capping_col:
           print(col)
           forecast[col] = capping_outlier(forecast, col)
```

```
{'forecast_price_energy_p2', 'forecast_cons_year', 'forecast_price_pow_p1',
'forecast_price_energy_p1', 'forecast_meter_rent_12m', 'forecast_cons_12m'}
forecast_price_energy_p2
forecast_cons_year
forecast_price_pow_p1
forecast_price_energy_p1
forecast_meter_rent_12m
forecast_cons_12m
```



```
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    self._setitem_single_column(loc, value, pi)
<ipython-input-37-792e83dc9994>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    forecast[col] = capping_outlier(forecast, col)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    self._setitem_single_column(loc, value, pi)
<ipython-input-37-792e83dc9994>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    forecast[col] = capping_outlier(forecast, col)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
<ipython-input-37-792e83dc9994>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
forecast[col] = capping_outlier(forecast, col)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
<ipython-input-37-792e83dc9994>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
forecast[col] = capping_outlier(forecast, col)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
<ipython-input-37-792e83dc9994>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
forecast[col] = capping_outlier(forecast, col)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
/home/brian/miniconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1720:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(loc, value, pi)
<ipython-input-37-792e83dc9994>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
forecast[col] = capping_outlier(forecast, col)

```
[38]: for col in forecast_col:
      print(col)
      detect_outlier(forecast, col)
```

forecast_cons_12m

```

Q1:513.23 Q3:2692.07750000000003
lower outlier :0.0, Upper outliers: 0.0
forecast_cons_year
Q1:0.0 Q3:1994.25
lower outlier :0.0, Upper outliers: 0.0
forecast_discount_energy
Q1:0.0 Q3:0.0
lower outlier :0.0, Upper outliers: 0.03597166998011928
forecast_meter_rent_12m
Q1:16.23 Q3:131.47
lower outlier :0.0, Upper outliers: 0.0
forecast_price_energy_p1
Q1:0.11523699999999999 Q3:0.146348
lower outlier :0.0, Upper outliers: 0.0
forecast_price_energy_p2
Q1:0.0 Q3:0.098837
lower outlier :0.0, Upper outliers: 0.0
forecast_price_pow_p1
Q1:40.606701 Q3:44.31137796
lower outlier :0.0, Upper outliers: 0.0

```

```
[39]: forecast.isnull().sum()
```

```

[39]: forecast_cons_12m          0
forecast_cons_year              0
forecast_discount_energy        0
forecast_meter_rent_12m         0
forecast_price_energy_p1        0
forecast_price_energy_p2        0
forecast_price_pow_p1           0
dtype: int64

```

2.4 Deal with categorical data

```

[40]: for col in cat_col:
        print(col)
        # Fill the categorical column with Mode
        train[col]
        mode = train[col].mode()[0]
        train[col].fillna(mode, inplace=True)

```

```

date_renewal
has_gas
date_modif_prod
date_activ
id
date_end

```

```
origin_up
channel_sales
```

```
[41]: train.isnull().sum()
```

```
[41]: id                0
      channel_sales    0
      cons_12m         0
      cons_gas_12m     0
      cons_last_month  0
      date_activ       0
      date_end         0
      date_modif_prod  0
      date_renewal     0
      forecast_cons_12m 0
      forecast_cons_year 0
      forecast_discount_energy 126
      forecast_meter_rent_12m 0
      forecast_price_energy_p1 126
      forecast_price_energy_p2 126
      forecast_price_pow_p1 126
      has_gas          0
      imp_cons         0
      margin_gross_pow_ele 13
      margin_net_pow_ele 13
      nb_prod_act      0
      net_margin       15
      num_years_antig  0
      origin_up        0
      pow_max          3
      churn            0
      dtype: int64
```

```
[42]: train[consumption_col] = consumption[consumption_col]
      train[forecast_col] = forecast[forecast_col]
```

```
[43]: train.isnull().sum()
```

```
[43]: id                0
      channel_sales    0
      cons_12m         0
      cons_gas_12m     0
      cons_last_month  0
      date_activ       0
      date_end         0
      date_modif_prod  0
      date_renewal     0
```

forecast_cons_12m	0
forecast_cons_year	0
forecast_discount_energy	0
forecast_meter_rent_12m	0
forecast_price_energy_p1	0
forecast_price_energy_p2	0
forecast_price_pow_p1	0
has_gas	0
imp_cons	0
margin_gross_pow_ele	13
margin_net_pow_ele	13
nb_prod_act	0
net_margin	15
num_years_antig	0
origin_up	0
pow_max	3
churn	0
dtype:	int64

```
[44]: for col in num_col:
      train[col].fillna(train[col].median(), inplace = True)
```

```
[45]: train.isnull().sum()
```

```
[45]: id                                0
      channel_sales                     0
      cons_12m                          0
      cons_gas_12m                      0
      cons_last_month                   0
      date_activ                        0
      date_end                          0
      date_modif_prod                   0
      date_renewal                      0
      forecast_cons_12m                 0
      forecast_cons_year                0
      forecast_discount_energy          0
      forecast_meter_rent_12m           0
      forecast_price_energy_p1          0
      forecast_price_energy_p2          0
      forecast_price_pow_p1             0
      has_gas                           0
      imp_cons                          0
      margin_gross_pow_ele              0
      margin_net_pow_ele                0
      nb_prod_act                       0
      net_margin                        0
      num_years_antig                   0
```

```
origin_up          0
pow_max            0
churn              0
dtype: int64
```

3 By far, we've already dealt with ourlier and fill the null values.
Next, use these dataset to build the model

```
[53]: from sklearn.preprocessing import LabelEncoder
      labelcoder = LabelEncoder()
```

```
[54]: for col in cat_col:
      train[col] = labelcoder.fit_transform(train[col])
```

```
[56]: train_data = train.drop(columns='churn',axis = 0)
      churn = train['churn']
```

```
[57]: from sklearn.model_selection import train_test_split
      # X_train,X_test,Y_train,Y_test = train_test_split(one_hot_ticket,Survived,
      #           ↪test_size=.30)
      X_train,X_test,Y_train,Y_test = train_test_split(train_data,churn, test_size=.
      ↪30)
```

```
[58]: # machine learning
      from sklearn.linear_model import LogisticRegression
      from sklearn.svm import SVC, LinearSVC
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.naive_bayes import GaussianNB
      from sklearn.linear_model import Perceptron
      from sklearn.linear_model import SGDClassifier
      from sklearn.tree import DecisionTreeClassifier
```

```
[59]: # Logistic Regression
      logreg = LogisticRegression()
      logreg.fit(X_train, Y_train)
      Y_pred = logreg.predict(X_test)
      acc_log = round(logreg.score(X_test, Y_test) * 100, 2)

      # Stochastic Gradient Descent
      sgd = SGDClassifier()
      sgd.fit(X_train, Y_train)
      Y_pred = sgd.predict(X_test)
      acc_sgd = round(sgd.score(X_test, Y_test) * 100, 2)
      acc_sgd
```

```

# Support Vector Machines
svc = SVC()
svc.fit(X_train, Y_train)
Y_pred = svc.predict(X_test)
acc_svc = round(svc.score(X_test, Y_test) * 100, 2)
acc_svc

# KNN
knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(X_train, Y_train)
Y_pred = knn.predict(X_test)
acc_knn = round(knn.score(X_test, Y_test) * 100, 2)
acc_knn

# Gaussian Naive Bayes
gaussian = GaussianNB()
gaussian.fit(X_train, Y_train)
Y_pred = gaussian.predict(X_test)
acc_gaussian = round(gaussian.score(X_test, Y_test) * 100, 2)
acc_gaussian

# Perceptron
perceptron = Perceptron()
perceptron.fit(X_train, Y_train)
Y_pred = perceptron.predict(X_test)
acc_perceptron = round(perceptron.score(X_test, Y_test) * 100, 2)
acc_perceptron

# Linear SVC
linear_svc = LinearSVC()
linear_svc.fit(X_train, Y_train)
Y_pred = linear_svc.predict(X_test)
acc_linear_svc = round(linear_svc.score(X_test, Y_test) * 100, 2)
acc_linear_svc

# Decision Tree
decision_tree = DecisionTreeClassifier()
decision_tree.fit(X_train, Y_train)
Y_pred = decision_tree.predict(X_test)
acc_decision_tree = round(decision_tree.score(X_test, Y_test) * 100, 2)
acc_decision_tree

# Random Forest

```



```

random_forest = RandomForestClassifier(n_estimators=100)
random_forest.fit(X_train, Y_train)
Y_pred = random_forest.predict(X_test)
random_forest.score(X_train, Y_train)
acc_random_forest = round(random_forest.score(X_test, Y_test) * 100, 2)
acc_random_forest

```

```

/home/brian/miniconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
/home/brian/miniconda3/lib/python3.8/site-packages/sklearn/svm/_base.py:985:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.

```

```
warnings.warn("Liblinear failed to converge, increase ")
```

[59]: 90.21

```

[60]: # drop string
models = pd.DataFrame({
    'Model': ['Support Vector Machines', 'KNN', 'Logistic Regression',
             'Random Forest', 'Naive Bayes', 'Perceptron',
             'Stochastic Gradient Decent', 'Linear SVC',
             'Decision Tree'],
    'Accuracy_Score': [acc_svc, acc_knn, acc_log,
                      acc_random_forest, acc_gaussian, acc_perceptron,
                      acc_sgd, acc_linear_svc, acc_decision_tree]})
models.sort_values(by='Accuracy_Score', ascending=False)

```

```

[60]:
      Model  Accuracy_Score
3  Random Forest          90.21
2  Logistic Regression      89.94
0  Support Vector Machines  89.92
7    Linear SVC            89.92
6  Stochastic Gradient Decent  89.56
1                KNN        87.74
4      Naive Bayes         86.91
8    Decision Tree         83.12
5      Perceptron         26.01

```

```
[61]: from sklearn.decomposition import PCA
      from sklearn.decomposition import TruncatedSVD
      from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
[62]: myPCA = PCA(10)
      mySVD = TruncatedSVD(10)
      myLDA = LinearDiscriminantAnalysis(10)
```

```
[64]: myPCA.fit(X_train)
```

```
[64]: PCA(n_components=10)
```

```
[65]: RX_train = myPCA.transform(X_train)
      RX_test = myPCA.transform(X_test)
```

```
[67]: # Random Forest
      random_forest = RandomForestClassifier(n_estimators=100)
      random_forest.fit(RX_train, Y_train)
      Y_pred = random_forest.predict(RX_test)
      random_forest.score(RX_train, Y_train)
      acc_random_forest = round(random_forest.score(RX_test, Y_test) * 100, 2)
      acc_random_forest
```

```
[67]: 90.14
```

```
[69]: # Logistic Regression
      logreg = LogisticRegression()
      logreg.fit(X_train, Y_train)
      Y_pred = logreg.predict(X_test)
      acc_log = round(logreg.score(X_test, Y_test) * 100, 2)
      acc_log
```

```
/home/brian/miniconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[69]: 89.94
```

```
[ ]:
```